

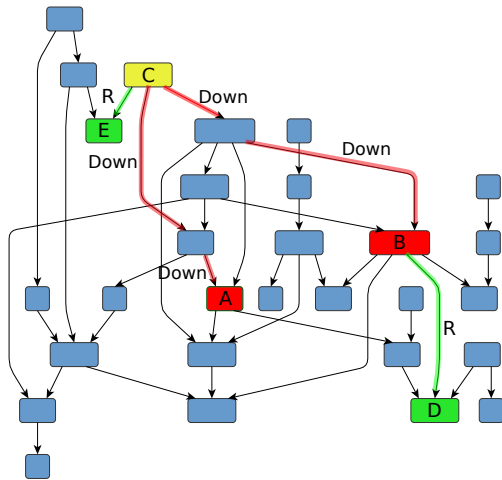
GLL-based Context-Free Path Querying for Neo4j

Vadim Abzalov, Vlada Pogozhelskaya, Vladimir Kutuev,
Olga Bachishche, **Semyon Grigorev**

Saint Petersburg State University

October 31, 2025


Formal Language Constrained Path Querying



Navigation through an edge-labeled graph

- **Path** specifies a **word** formed by the labels of the edges
- **Paths constraint** is a **language**: the word specified by the path should be in the given language
- The expressiveness of constraints is related to **formal languages classes**


Context-Free Path Queries (CFPQ)



Context-free languages as constraints

- Are nodes A and B on the same level of hierarchy?
- Is there a path of form $\overline{\text{Down}}^n \text{Down}^n$ between A and B?
- Context-free grammar:
 $\text{SameLvl} \rightarrow \overline{\text{Down}} \text{SameLvl} \text{Down} \mid \varepsilon$

Context-Free Path Queries (CFPQ)



Context-free languages as constraints

- Are nodes A and B on the same level of hierarchy?
- Is there a path of form $\overline{\text{Down}}^n \text{Down}^n$ between A and B?
- Context-free grammar:
 $\text{SameLvl} \rightarrow \overline{\text{Down}} \text{SameLvl} \text{Down} \mid \varepsilon$

Applications

- Static code analysis [T. Reps, et al, 1995]
- Graph segmentation [H. Miao, et al, 2019]
- Bio data analysis [P. Sevon, et al, 2008]
- ...

Problem Statement

- J. Kuijpers, et al¹: existing algorithms are too slow to be used in practical applications (in the context of Neo4j)
- Reachability in the focus
 - ▶ Paths needed in some applications
 - ▶ Not for all pairs, but for specified start vertices

? How to create faster multiple source context-free all paths querying algorithm?

¹Jochem Kuijpers, George Fletcher, Nikolay Yakovets, and Tobias Lindaaker. 2019. An Experimental Study of Context-Free Path Query Evaluation Methods.

- **Generalized LL (GLL)²** as a base
 - ▶ Arbitrary grammars (including left-recursive and ambiguous) without transformations
 - ▶ **Shared Packed Parse Forest (SPPF)** is a native representation of all paths
 - ▶ Directed — native support of source vertices
- **Recursive State Machine (RSM)** to represent constraints
 - ▶ Instead of grammar in (E)BNF³

²A. Afroozeh, Anastasia Izmaylova. Faster, Practical GLL Parsing. 2015

³Right part of the rule is a regular expression over terminals and nonterminals

Definitions


- **Descriptor** $d = (u, q, g)$ — configuration of recognizer/parser
 - ▶ u — vertex of the input graph
 - ▶ q — state of the RSM (constraint)
 - ▶ g — current top of the stack (vertex of GSS)
- **Graph Structured Stack (GSS)** — compact⁴ representation of multiple stacks
 - ▶ Vertex (M, u) — Handling of nonterminal M is started from vertex u of the input graph
 - ▶ Each edge is labelled with p — state of the RSM to continue from after pop
- **Shared Packed Parse Forest (SPPF)** — compact⁵ representation of all possible derivation trees
 - ▶ **Terminal** and **nonterminal** nodes
 - ▶ **Packed, intermediate, etc** — special nodes to reuse subtrees
- **Recursive State Machine (RSM)** — finite-automata-like representation of context-free language
 - ▶ DFA representation of right part of rule for each nonterminal

⁴Cubic in worst case

⁵Cubic in worst case

An Example


Input graph:



Let 1 be a start vertex


Grammar (query, constraint): $S \rightarrow a S b \mid a b$

RSM:



Generalized LL for CFPQ: The Idea


Current descriptor: (u, a, g)



Just read the terminal


New descriptor: (v, p, g)
// For all terminal edges

Call: start handling
of M on position u



New descriptor: (u, r, h)
// r : start state for M
// For all Nonterminal edges

a is a final state



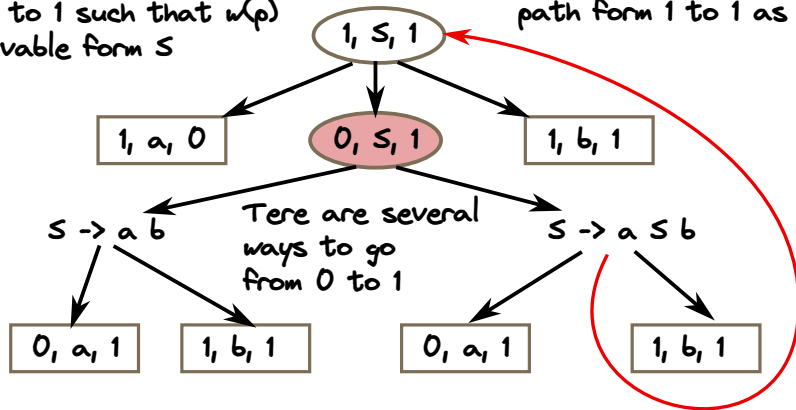
Pop is not destructive.
Just move pointer
alongside outgoing edges

New descriptors: (u, p, h)
 (u, r, f)
// For each outgoing edge

SPPF is a Representation of All Paths of Interest


There is at least one path from 1 to 1 such that $w(p)$ is derivable from S

Cyclic references:
path from 0 to 1 contains
path from 1 to 1 as a subpath



Trees And Paths

without cycle




Trees extracted from SPPF for the following paths:

$1 \xrightarrow{a} 0 \xrightarrow{a} 1 \xrightarrow{b} 1 \xrightarrow{b} 1$


$1 \xrightarrow{a} 0 \xrightarrow{a} 1 \xrightarrow{a} 0 \xrightarrow{a} 1 \xrightarrow{b} 1 \xrightarrow{b} 1 \xrightarrow{b} 1 \xrightarrow{b} 1$

cycle unrolled once




Context-Free Languages Are Closed Under Intersection With Regular Ones

Input graph: regular language




Query: context-free language



Result of intersection

SPPF



Respective context-free grammar

$(1, S, 1) \rightarrow (1, a, 0) (0, S, 1) (1, b, 1)$
 $(0, S, 1) \rightarrow (0, a, 1) (1, b, 1)$
 $(0, S, 1) \rightarrow (0, a, 1) (1, S, 1) (1, b, 1)$

Nonterminals

Terminals
(edges)

- Generic CFPQ solver (graph is abstraction)
- Neo4j as graph storage (no Cypher extension to support CFPQ)

Evaluation Setup

- Ubuntu 20.04, Intel Core i7-6700 CPU, 3.4GHz, DDR4 64Gb RAM
- Neo4j 5.12.0
 - ▶ Unlimited memory usage per transaction
- JVM was configured to use 55Gb

Graph name	V	E	#subClassOf	#type
Core	1 323	2 752	178	0
Pathways	6 238	12 363	3 117	3 118
Go_hierarchy	45 007	490 109	490 109	0
Enzyme	48 815	86 543	8 163	14 989
Eclass	239 111	360 248	90 962	72 517
Geospecies	450 609	2 201 532	0	89 065
Go	582 929	1 437 437	94 514	226 481
Taxonomy	5 728 398	14 922 125	2 112 637	2 508 635

Queries:

$$Q_1 : \quad S \rightarrow \overline{\text{subClassOf}} \ S \ \text{subClassOf} \mid \overline{\text{type}} \ S \ \text{type} \\ \mid \overline{\text{subClassOf}} \ \text{subClassOf} \mid \overline{\text{type}} \ \text{type}$$


$$Q_2 : \quad S \rightarrow \overline{\text{subClassOf}} \ S \ \text{subClassOf} \mid \text{subClassOf}$$

$$\text{reg}_1 : \quad S \rightarrow (\text{subClassOf} \mid \text{type})^*$$


$$\text{reg}_2 : \quad S \rightarrow \text{subClassOf}^* \cdot \text{type}^*$$

CFPQ reachability speedup (RSM over CFG) on RDF graphs

Q_1




Q_2




CFPQ results (Neo4j vs RedisGraph)

Q_1




Q_2




RPQ results for Neo4j⁶

reg₁



reg₂



⁶Native solution failed with OOM on last two graphs

Conclusion And Future Research

✓ GLL-based CFPQ algorithm

- ▶ RSM is promising representation of CFLs in context of CFPQ
- ▶ In some cases faster than linear-algebra-based approach
- ▶ Comparable with native RPQ (for Neo4j)

⚙️ Implementation of paths extraction strategies



Parallel version of GLL

- 💡 All descriptors can be handled independently
- ☹️ Complex global shared structures (GSS, SPPF, etc) — synchronization needed



Incremental version of GLL