



# Brahma.FSharp: Power of Functional Programming to Create Portable GPGPU-enabled .NET Applications

Nikolai Ponomarev, Vladimir Kutuev, Semyon Grigorev

St. Petersburg State University

October 7, 2025

# General-Purpose Computing on GPUs

- Not only HPC and scientific computations
- Widely used in regular business applications
  - ▶ Images and video processing
  - ▶ AI/ML
  - ▶ Financial analysis
  - ▶ ...
- Highest performance vs development simplification
  - ▶ Native execution vs managed environments (JVM, .NET, etc)
  - ▶ Low-level languages vs high-level ones
  - ▶ ...
- Functional languages for GPGPU
  - ▶ Type-safety, flexibility, advanced optimizations
  - ▶ Futhark, Lift, AnyDSL, ...

- .NET — one of popular platforms for business applications development
- F# — functional-first programming language for .NET
  - ▶ Statically strongly typed
    - ★ Type-safe development of generic kernels
  - ▶ F# Code Quotations
    - ★ Compile-time metaprogramming
    - ★ Runtime access to full-featured AST
    - ★ Statically typed
    - ★ Type-safe composition
  - ▶ F# Mailbox Processor
    - ★ Integrated language-native primitive for message-passing-based asynchronous programming

# F# for GPGPU Programming: Design Decision Points

- **Target technology**
  - ▶ CUDA (e.g. *Alea GPU*) — mature stable infrastructure but Nvidia only
  - ▶ OpenCL (e.g. *FSCL*) — wide range of devices and vendors
- **Translator input**
  - ▶ Bytecode (e.g. *ILGPU*) — one tool for all languages on the platform
  - ▶ Source code (e.g. *FSCL*) — an ability to use language-specific features
- **Control level**
  - ▶ Fine-grained (e.g. *ILGPU*)
  - ▶ black-box-like automatization (e.g. *FSCL*)
- **Type safety**
  - ▶ None
  - ▶ Manual type annotations
  - ▶ Automatic types inference and checking
- **Focus**
  - ▶ Alea GPU: HPC, AI/ML, bindings for existing Cuda libs
  - ▶ FSCL: automatic computation scheduling in heterogeneous systems
  - ▶ ...

is a **F# quotation to OpenCL translator** and respective **runtime** to utilize GPGPUs in F# applications

- **Target technology:** OpenCL
- **Translator input:** source code (using quotations)
  - ▶ Not only primitive types, but also discriminated unions, structs, records
  - ▶ Pattern matching, mutable and immutable bindings, nested bindings
- **Control level:** fine-grained memory management and kernels compilation process
- **Type safety:** automatic
  - ▶ Ten lines of types' manipulation magic that carefully build type for compiled kernel<sup>1</sup>
- **Focus:** safety, static checks, development simplification
- Mailbox Processor wrapper for command queue
  - ▶ Native way to integrate communication with GPGPUs into asynchronous pipelines

---

<sup>1</sup>... and brakes some F# compiler optimizations

# Research Questions

- Can we create code portable across different devices and vendors?
- Can we create generic type-safe kernels?
- Can we utilize well-established optimizations if necessary?
- Does Mailbox Processor allows one to create asynchronous workflows that utilize heterogeneous devices?

## An Example of Generic Matrix Multiplication Kernel

```
let mXmKernel (opAdd: Quotations.Expr<'a -> 'b -> 'a>)  
               (opMult: Quotations.Expr<'e -> 'f -> 'b>)  
               (zero: Quotations.Expr<'a>) ... (* other parameters *) =  
    ... // Supplementary code  
let kernel = <@ fun 2dRange m1 m2 res -> // Quoted code  
    ...  
    let acc = %zero // Embedded identity value  
    let lBuf = localArray lws // captured from context  
    ...  
    acc <- (%opAdd) acc ((%opMult) x y) // Embedded operations  
    ... @>  
    ... // Supplementary code  
  
let intArithmeticKernel = mXmKernel <@ (+) @> <@ ( * ) @> <@ 0 @>  
let intMinPlusKernel = mXmKernel <@ (min) @> <@ (+) @> <@ Int.MaxValue @>
```

## An Example of Generic Matrix Multiplication Kernel

```
let mXmKernel (opAdd: Quotations.Expr<'a -> 'b -> 'a>)  
              (opMult: Quotations.Expr<'e -> 'f -> 'b>)  
              (zero: Quotations.Expr<'a>) ... (* other parameters *) =  
    ... // Supplementary code  
let kernel = <@ fun 2dRange m1 m2 res -> // Quoted code  
    ...  
    let acc = %zero // Embedded identity value  
    let lBuf = localArray lws // captured from context  
    ...  
    acc <- (%opAdd) acc ((%opMult) x y) // Embedded operations  
    ... @>  
    ... // Supplementary code  
  
let intArithmeticKernel = mXmKernel <@ (+) @> <@ ( * ) @> <@ 0 @>  
let intMinPlusKernel = mXmKernel <@ (min) @> <@ (+) @> <@ Int.MaxValue @>
```



## An Example of Generic Matrix Multiplication Kernel

```
let mXmKernel (opAdd: Quotations.Expr<'a -> 'b -> 'a>)  
               (opMult: Quotations.Expr<'e -> 'f -> 'b>)  
               (zero: Quotations.Expr<'a>) ... (* other parameters *) =  
    ... // Supplementary code  
    let kernel = <@ fun 2dRange m1 m2 res -> // Quoted code  
        ...  
        let acc = %zero // Embedded identity value  
        let lBuf = localArray lws // captured from context  
        ...  
        acc <- (%opAdd) acc ((%opMult) x y) // Embedded operations  
        ... @>  
    ... // Supplementary code  
  
let intArithmeticKernel = mXmKernel <@ (+) @> <@ ( * ) @> <@ 0 @>  
let intMinPlusKernel = mXmKernel <@ (min) @> <@ (+) @> <@ Int.MaxValue @>
```

## An Example of Generic Matrix Multiplication Kernel

```
let mXmKernel (opAdd: Quotations.Expr<'a -> 'b -> 'a>)  
              (opMult: Quotations.Expr<'e -> 'f -> 'b>)  
              (zero: Quotations.Expr<'a>) ... (* other parameters *) =  
... // Supplementary code  
let kernel = <@ fun 2dRange m1 m2 res -> // Quoted code  
...  
let acc = %zero // Embedded identity value  
let lBuf = localArray lws // captured from context  
...  
acc <- (%opAdd) acc ((%opMult) x y) // Embedded operations  
... @>  
... // Supplementary code  
  
let intArithmeticKernel = mXmKernel <@ (+) @> <@ ( * ) @> <@ 0 @>  
let intMinPlusKernel = mXmKernel <@ (min) @> <@ (+) @> <@ Int.MaxValue @>
```

## An Example of Generic Matrix Multiplication Kernel

```
let mXmKernel (opAdd: Quotations.Expr<'a -> 'b -> 'a>)  
              (opMult: Quotations.Expr<'e -> 'f -> 'b>)  
              (zero: Quotations.Expr<'a>) ... (* other parameters *) =  
    ... // Supplementary code  
let kernel = <@ fun 2dRange m1 m2 res -> // Quoted code  
    ...  
    let acc = %zero // Embedded identity value  
    let lBuf = localArray lws // captured from context  
    ...  
    acc <- (%opAdd) acc ((%opMult) x y) // Embedded operations  
    ... @>  
    ... // Supplementary code  
  
let intArithmeticKernel = mXmKernel <@ (+) @> <@ ( * ) @> <@ 0 @>  
let intMinPlusKernel = mXmKernel <@ (min) @> <@ (+) @> <@ Int.MaxValue @>
```

## An Example of Generic Matrix Multiplication Kernel

```
let mXmKernel (opAdd: Quotations.Expr<'a -> 'b -> 'a>)  
               (opMult: Quotations.Expr<'e -> 'f -> 'b>)  
               (zero: Quotations.Expr<'a>) ... (* other parameters *) =  
... // Supplementary code  
let kernel = <@ fun 2dRange m1 m2 res -> // Quoted code  
...  
let acc = %zero // Embedded identity value  
let lBuf = localArray lws // captured from context  
...  
acc <- (%opAdd) acc ((%opMult) x y) // Embedded operations  
... @>  
... // Supplementary code  
  
let intArithmeticKernel = mXmKernel <@ (+) @> <@ ( * ) @> <@ 0 @>  
let intMinPlusKernel = mXmKernel <@ (min) @> <@ (+) @> <@ Int.MaxValue @>
```

## An Example of Generic Matrix Multiplication Kernel

```
let mXmKernel (opAdd: Quotations.Expr<'a -> 'b -> 'a>)  
              (opMult: Quotations.Expr<'e -> 'f -> 'b>)  
              (zero: Quotations.Expr<'a>) ... (* other parameters *) =  
    ... // Supplementary code  
let kernel = <@ fun 2dRange m1 m2 res -> // Quoted code  
    ...  
    let acc = %zero // Embedded identity value  
    let lBuf = localArray lws // captured from context  
    ...  
    acc <- (%opAdd) acc ((%opMult) x y) // Embedded operations  
    ... @>  
    ... // Supplementary code  
  
let intArithmeticKernel = mXmKernel <@ (+) @> <@ ( * ) @> <@ 0 @>  
let intMinPlusKernel = mXmKernel <@ (min) @> <@ (+) @> <@ Int.MaxValue @>
```



# Matrix Multiplication Evaluation: Environment

- Optimizations
  - ▶ Tiling in local and private memory<sup>2</sup>
  - ▶ Square matrices and square tiles
- Hardware
  - ▶ **Lenovo**: Intel Core i7-8550U CPU, NVIDIA GeForce MX150 and Intel UHD Graphics 620 GPUs
  - ▶ **Zen**: Intel Core i5-1340P CPU, Intel Iris Xe Graphics G7 80EUs GPU
  - ▶ **MILK-V**: SpacemiT M1 CPU, IMG BXE-2-32 GPU
- Software
  - ▶ **CLBlast**: optimized OpenCL-based BLAS implementation tuned for low-power mobile GPUs
  - ▶ **OpenBLAS**: optimized CPU-based BLAS implementation
- **PoCL** to execute OpenCL-based solutions on CPUs

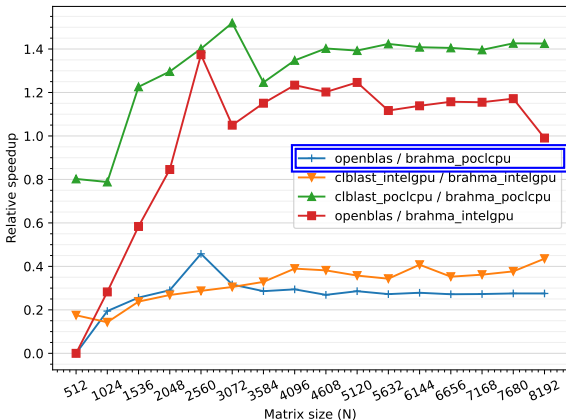

---

<sup>2</sup>Inspired by “[Tutorial: OpenCL SGEMM tuning for Kepler](#)” by Cedric Nugteren

# Matrix Multiplication Performance: **Lenovo** and **Zen**





# Matrix Multiplication Performance: **Lenovo** and **Zen**









# Matrix Multiplication Performance: **Lenovo** and **Zen**




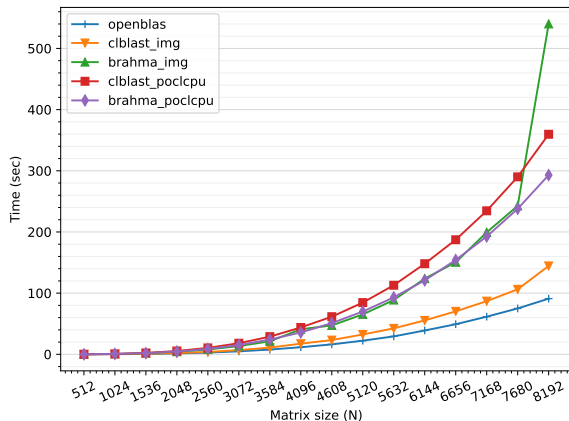
# Matrix Multiplication Performance: **Lenovo** and **Zen**



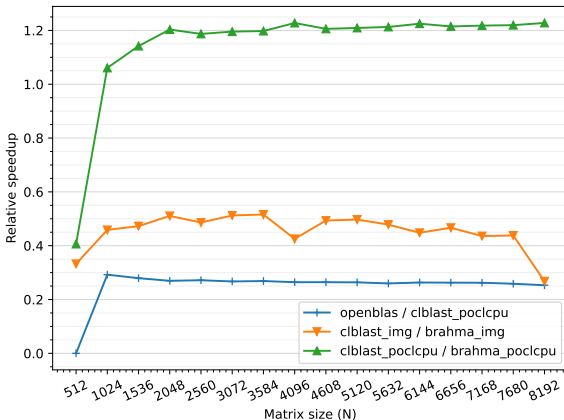
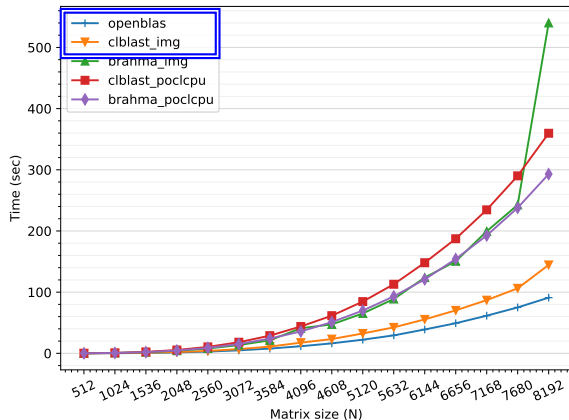
# Matrix Multiplication Performance: **Lenovo** and **Zen**



# Matrix Multiplication Performance: MILK-V



# Matrix Multiplication Performance: MILK-V



## Image Processing Using Mailbox Processor<sup>3</sup>

```
let imgProcessor filter (imgSaver: MailboxProcessor<_>) =  
    MailboxProcessor.Start(fun inbox ->  
        let rec loop ... = async { // Async message processing loop  
            let! msg = inbox.Receive() // Load message  
            match msg with  
            | EOS ch -> // Handle end of stream  
                imgSaver.PostAndReply EOS  
                ch.Reply()  
            | Img img -> // Handle image  
                let filtered = filter img // Convolution  
                imgSaver.Post (Img filtered)  
                return! loop ... } // Got to next message  
        loop ...)
```

---

<sup>3</sup>Basic version

## Image Processing Using Mailbox Processor<sup>3</sup>

```
let imgProcessor filter (imgSaver: MailboxProcessor<_>) =  
    MailboxProcessor.Start(fun inbox ->  
        let rec loop ... = async { // Async message processing loop  
            let! msg = inbox.Receive() // Load message  
            match msg with  
            | EOS ch -> // Handle end of stream  
                imgSaver.PostAndReply EOS  
                ch.Reply()  
            | Img img -> // Handle image  
                let filtered = filter img // Convolution  
                imgSaver.Post (Img filtered)  
                return! loop ... } // Got to next message  
        loop ...)
```

---

<sup>3</sup>Basic version

## Image Processing Using Mailbox Processor<sup>3</sup>

```
let imgProcessor filter (imgSaver: MailboxProcessor<_>) =  
    MailboxProcessor.Start(fun inbox ->  
        let rec loop ... = async { // Async message processing loop  
            let! msg = inbox.Receive() // Load message  
            match msg with  
            | EOS ch -> // Handle end of stream  
                imgSaver.PostAndReply EOS  
                ch.Reply()  
            | Img img -> // Handle image  
                let filtered = filter img // Convolution  
                imgSaver.Post (Img filtered)  
                return! loop ... } // Got to next message  
        loop ...)
```

---

<sup>3</sup>Basic version



## Image Processing Using Mailbox Processor<sup>3</sup>

```
let imgProcessor filter (imgSaver: MailboxProcessor<_>) =  
    MailboxProcessor.Start(fun inbox ->  
        let rec loop ... = async { // Async message processing loop  
            let! msg = inbox.Receive() // Load message  
            match msg with  
            | EOS ch -> // Handle end of stream  
                imgSaver.PostAndReply EOS  
                ch.Reply()  
            | Img img -> // Handle image  
                let filtered = filter img // Convolution  
                imgSaver.Post (Img filtered)  
                return! loop ... } // Got to next message  
        loop ...)
```

---

<sup>3</sup>Basic version

## Image Processing Using Mailbox Processor<sup>3</sup>

```
let imgProcessor filter (imgSaver: MailboxProcessor<_>) =  
    MailboxProcessor.Start(fun inbox ->  
        let rec loop ... = async { // Async message processing loop  
            let! msg = inbox.Receive() // Load message  
            match msg with  
            | EOS ch -> // Handle end of stream  
                imgSaver.PostAndReply EOS  
                ch.Reply()  
            | Img img -> // Handle image  
                let filtered = filter img // Convolution  
                imgSaver.Post (Img filtered)  
                return! loop ... } // Got to next message  
        loop ...)
```

---

<sup>3</sup>Basic version

## Image Processing Using Mailbox Processor<sup>3</sup>

```
let imgProcessor filter (imgSaver: MailboxProcessor<_>) =  
    MailboxProcessor.Start(fun inbox ->  
        let rec loop ... = async { // Async message processing loop  
            let! msg = inbox.Receive() // Load message  
            match msg with  
            | EOS ch -> // Handle end of stream  
                imgSaver.PostAndReply EOS  
                ch.Reply()  
            | Img img -> // Handle image  
                let filtered = filter img // Convolution  
                imgSaver.Post (Img filtered)  
                return! loop ... } // Got to next message  
        loop ...)
```

---

<sup>3</sup>Basic version

## Image Processing Using Mailbox Processor<sup>3</sup>

```
let imgProcessor filter (imgSaver: MailboxProcessor<_>) =  
    MailboxProcessor.Start(fun inbox ->  
        let rec loop ... = async { // Async message processing loop  
            let! msg = inbox.Receive() // Load message  
            match msg with  
            | EOS ch -> // Handle end of stream  
                imgSaver.PostAndReply EOS  
                ch.Reply()  
            | Img img -> // Handle image  
                let filtered = filter img // Convolution  
                imgSaver.Post (Img filtered)  
                return! loop ... } // Got to next message  
        loop ...)
```

---

<sup>3</sup>Basic version

# Evaluaiton of Image Processing

- Environment

- ▶ **Lenovo**: Intel Core i7-8550U CPU, NVIDIA GeForce MX150 and Intel UHD Graphics 620 GPUs
- ▶ Stream processing: 420 images (1Gb of data), load–process–store
- ▶  $5 \times 5$  kernels sequence: 3 Gaussian blur, 1 edge detection

- Results

- ▶ 64 seconds on Nvidia GPU
- ▶ 97 seconds on Intel GPU
- ▶ 40 seconds using Intel and Nvidia simultaneously: **30% speedup**

# Conclusion

- Brahma.FSharp allows one to create portable solutions
  - ▶ We want to create homogeneous code<sup>4</sup> to simplify development and portability process<sup>5</sup>
- Brahma.FSharp allows one to create highly-parameterizable type-safe kernels
  - ▶ Code support simplification: single highly-optimized and type-safe kernel<sup>6</sup>
- Brahma.FSharp allows one to utilize well-established optimizations
  - ▶ Including manipulations with local and private memory
- Mailbox Processor allows one
  - ▶ To create asynchronous workflows that utilize heterogeneous environment in F#-native way
  - ▶ To utilize heterogeneous devices to increase performance
- We should reduce runtime overhead
- In case of matrix multiplication, there is a room for kernels tuning

---

<sup>4</sup>Even .NET-only, without wrappers for native libraries

<sup>5</sup>Write once, run everywhere. In some cases.

<sup>6</sup>But without type-specific optimizations

# Useful Links

- Brahma.FSharp sources:  
<https://github.com/YaccConstructor/Brahma.FSharp>
- Environment for evaluation and examples:  
<https://github.com/gsvgit/ImageProcessing>



- Email:  
s.v.grigoriev@mail.spbu.ru
- GitHub: [gsvgit](#)
- Google Scholar: [Semyon Grigorev](#)
- DBLP: [Semyon V. Grigorev](#)