

Graph Analysis via Graph Rewriting: Sparse Linear Algebra Based Graph Analysis in Interaction Nets

Danil Zaripov^{1,*}, Efim Kubyshkin¹, Nikolai Ponomarev¹, Georgy Sichkar¹, Vladimir Zaikin¹ and Semyon Grigorev¹

¹*Saint Petersburg State University, 7-9 Universitetskaya Embankment, St. Petersburg, 199034, Russia*

Abstract

A clear and well-documented $\text{\LaTeX}$ document is presented as an article formatted for publication by CEUR-WS in a conference proceedings. Based on the “ceurart” document class, this article presents and explains many of the common variations, as well as many of the formatting elements an author may use in the preparation of the documentation of their work.

Keywords

Interaction Nets, graph rewriting, graph analysis algorithms, GraphBLAS, sparse linear algebra, quad-tree

1. Introduction

Graph analysis has become fundamental to modern data science, powering applications ranging from recommendation systems in social networks and fraud detection in financial transactions to protein interaction analysis in bioinformatics and static code analysis. As graph sizes grow to billions of vertices and edges, the need for efficient, scalable implementations becomes critical. However, graph analysis algorithms are difficult to parallelize due to their irregular memory access patterns, poor data locality, and load imbalance across computational units.

The GraphBLAS standard [1] addresses these challenges by providing a foundation for constructing high-performance parallel graph analysis algorithms. By expressing graph algorithms in the language of linear algebra — using generic sparse vectors and matrices as core data structures, together with a rich set of highly parameterizable operations (matrix multiplication, element-wise operations, etc.) — GraphBLAS offers a high-level, mathematically rigorous framework for algorithm construction. Operations over sparse generic matrices and vectors parameterized by semiring-like structures provide a powerful basis for expressing a wide range of graph analysis algorithms [2].

Despite its successes, GraphBLAS and its implementations face certain challenges including the implementation of high-performance generic operations over sparse data structures. Imbalanced subtasks together with irregular data access patterns make load balancing — and thus efficient parallel implementation — hard. Moreover, graph analysis algorithms are not single operations but nontrivial chains of operations over matrices and vectors. This leads to the problem of intermediate data structures. Kernel fusion could solve these problems, but it is too difficult to implement automatic optimizations for sparse generic computations, so it is done manually in the most frequent cases.

At the same time, interaction nets [3] — a graph-rewriting model introduced by Yves Lafont — offer a computational model that is inherently well-suited for irregular parallelism. In interaction nets, computation proceeds through local rewriting rules applied to pairs of connected nodes (so-called active pairs). Crucially, multiple non-interfering active pairs can be reduced simultaneously, making interaction nets a promising foundation for parallel irregular computations, including graph analysis.

Woodstock’22: Symposium on the irreproducible science, June 07–11, 2022, Woodstock, NY

*Corresponding author.

✉ st129060@student.spbu.ru (D. Zaripov); e.kubyshkin@spbu.ru (E. Kubyshkin); n.ponomarev@spbu.ru (N. Ponomarev); georgy.sichkar@yandex.ru (G. Sichkar); st119057@student.spbu.ru (V. Zaikin); s.v.grigoriev@mail.spbu.ru (S. Grigorev)

🆔 0009-0003-1086-9279 (D. Zaripov); 0009-0009-5531-1758 (E. Kubyshkin); 0009-0000-2382-5687 (N. Ponomarev); 0009-0005-4684-9610 (G. Sichkar); 0009-0003-3417-990X (V. Zaikin); 0000-0002-7966-0698 (S. Grigorev)



© 2022 Copyright for this paper by its authors. Use permitted under Creative Commons License Attribution 4.0 International (CC BY 4.0).

Furthermore, interaction nets are relatively close to functional programming languages, making it easy to transfer design approaches and ideas from the functional world.

However, comparisons between existing interaction-net-based evaluators remain very limited and are often performed only on basic benchmarks such as Fibonacci, Ackermann, sorting, and list traversal.

We believe that interaction nets can provide a powerful foundation for implementing graph analysis algorithms. The long-term goal of our work is to investigate how interaction nets can help solve challenges with irregular parallelism and, possibly, with kernel fusion. In this work, we do not focus on the GraphBLAS API implementation, but we are inspired by it: we implemented a minimal set of generic data structures and operations sufficient to implement several classic graph analysis algorithms.

This work is in progress, and our current contributions are as follows.

1. We implemented a minimal, functionally equivalent subset of the functions over tree-represented matrices and vectors in F# (a mature functional-first language, chosen for ease of prototyping and testing), Vine, and Inpla. On this foundation, we implemented the following graph analysis algorithms: BFS, SSSP, and triangle counting. Thus, we enable direct comparison of the respective execution models and runtime systems.
2. We evaluated the implemented algorithms on a set of real-world and synthetic graphs and highlight directions for future research.


2. Background

In this section, we introduce tree-based matrix representation, provide the basics of interaction nets, and motivate our choice of interaction net implementations used in our evaluation.

2.1. Tree-based Representation for Sparse Matrices and Vectors

To store sparse matrices and vectors efficiently, we employ a quad-tree data structure. The matrix is recursively subdivided into four quadrants — northwest (NW), northeast (NE), southwest (SW), and southeast (SE) — until each block becomes either uniform or empty. Figure 1 shows an example of a matrix in the quad-tree form.

Figure 1: Example of the matrix and its quad-tree representation



This recursive partitioning allows large uniform blocks to be represented implicitly, thereby reducing the memory required to store the matrix. Unlike flat array formats such as CSR or CSC, which are standard in libraries like SuiteSparse:GraphBLAS, the quad-tree naturally supports a divide-and-conquer execution model and adapts dynamically to irregular sparsity patterns. This structure simplifies parallelization of element-wise, matrix-matrix, and matrix-vector operations, as independent sub-blocks can be processed recursively without global synchronization.

2.2. Interaction Nets

Interaction Nets is a graph-rewriting computation model proposed by Y. Lafont in 1989 [3] and can be defined as follows.

Definition 1 (Agent). An agent is a symbol from a fixed alphabet Σ , for which:

- an arity is defined — a function $Ar : \Sigma \rightarrow \mathbb{N} \cup \{0\}$;
- there is one principal port and a number of auxiliary ports equal to its arity.

Definition 2 (Net). A net over the alphabet Σ is a labeled undirected multi-graph in which:

- each node is labeled by an agent;
- edges can only exist between ports;
- at most one edge enters each port.

Let $\mathcal{N}_\Sigma$ denote the set of all possible nets over the alphabet Σ .

We say that port p_a is connected to port p_b if there is an edge between them. We say that a port is *free* if no edge originates from it. A net can also consist simply of a set of edges, in which case each end of an edge is considered a free port.

Definition 3 (Interface). The interface $I(n)$ is the set of all free ports of a net n .


Computation in Interaction Nets is defined in terms of *reductions*: rewriting an *active pair* into a net as defined in *reduction rules*. An active pair consists of two agents connected by an edge between their principal ports.

Definition 4 (Reduction Rule). A reduction rule is a graph rewriting rule. Let R be the set of reduction rules. More formally, a reduction rule is a triplet $r \in \Sigma \times \Sigma \times \mathcal{N}_\Sigma$, for which the following properties hold:

- **Symmetry:** $(l_1, l_2, n) \in R \Rightarrow (l_2, l_1, n) \in R$
- **Determinism:** $(l_1, l_2, n_1) \in R \wedge (m_1, m_2, n_2) \in R \Rightarrow (l_1, l_2) \neq (m_1, m_2)$
- **Interface Preservation:** $\forall (l_1, l_2, n) \in R, Ar(l_1) + Ar(l_2) = \|I(n)\|$

An example of the computation of the net encoding of the term $(0 + 1) + (0 + 1)$ is shown in Figure 2. The set of agents is $\Sigma = \{0, S, Add\}$ with arities $Ar(0) = 0$, $Ar(S) = 1$ and $Ar(Add) = 2$. Reduction rules shown in Figure 3. The left auxiliary ports of the active *Add*-agents are connected to the net representing 1, encoded in Peano numerals as $S(0)$, and the right auxiliary ports are connected to the third *Add*-agent. There is one free port — one of the auxiliary ports of the *Add*-agent. This constitutes the interface.

Figure 2: Example of computation in interaction nets: reduction of the term $(1 + 0) + (1 + 0)$ encoded in Peano numerals



As shown in Lafont's work [3], λ -calculus can be expressed through this formalism that makes interaction nets general purpose computation model. On the other hand, agents and reduction rules are not predefined, so one can define custom sets of agents and reduction rules, creating domain-specific (or problem-specific) models. Additionally, reductions can be performed independently, and it has been proven that their order does not matter. This is the primary reason for the natural suitability of Interaction Nets for parallel computation. These properties make Interaction Nets a promising platform for graph analysis tasks, as parallel graph analysis has an irregular nature. At the same time, if it is possible to introduce a powerful language for describing graph analysis algorithms (e.g., linear algebra), then accelerating the final solution may be possible by creating specific agents and rules.




Figure 3: Reduction rules for Peano numerals

2.3. Implementations of Interaction Nets

Currently, we focus on Interaction Nets implementations that provide their own runtime. Consequently, we omit solutions that embed Interaction Nets into other languages (e.g., OCaml [4] or Haskell) or those that merely provide an interpreter. To the best of our knowledge, three such projects are actively developed: Inpla¹ [5], Vine², and HVM³.

Although the HVM project is promising — offering a high-level language (Bend) that compiles to an interaction-net-based intermediate representation and an optional GPGPU-accelerated runtime — it is not yet stable enough for our purposes⁴. Therefore, we selected Inpla and Vine for evaluation.

Inpla provides a low-level interaction nets programming language with a parallel runtime. Programming in Inpla requires defining a custom set of rewriting rules for a chosen set of agents. The Inpla program for the interaction net from Figure 2 is shown in Listing 1. Computation runs on its own parallel virtual machine, implemented in C.

```
(* Defining custom rules *)
Add(r, x) >< Z => r ~ x; (* Rule presented in Figure 3a *)
Add(r, x) >< S(y) => Add(r, S(x)) ~ y; (* Rule presented in Figure 3b *)
(* (1 + 0) + (1 + 0): two active pairs (Fig 2) *)
Add(x, S(Z)) ~ Z, Add(Add(r, x), S(Z)) ~ Z;
(* Printing the result and freeing the `r` name *)
r; free r;
```

Listing 1: Peano numerals addition in Inpla

Vine, in contrast to Inpla, is a high-level Rust-like language with a rich standard library (also implemented in Vine). It compiles to a low-level interaction-net-based language called IVY, which consists of a fixed set of agents and reduction rules — specifically, a slightly extended version of the standard basis for lambda-calculus. The compiled program then runs on a parallel virtual machine implemented in Rust.

3. Implementation Details

In this section we describe our implementations: in F# (used as a prototype and reference for other implementations), in Inpla (a low-level language), and in Vine (a high-level language). We present the data structures and functions, list graph analysis algorithms (with short descriptions), and discuss details related to the implementation of the described structures and functions in selected languages.

¹Inpla project: <https://github.com/inpla/inpla>

²Vine project: <https://github.com/VineLang/vine>.

³HVM project: <https://github.com/HigherOrderCO/HVM2>.

⁴For example, open issues remain regarding failed tests (<https://github.com/HigherOrderCO/HVM2/issues/437>) and the failure to accelerate quad-tree-based computations (<https://github.com/HigherOrderCO/Bend/issues/738>).

```

type treeValue<'value> = | Dummy | UserValue of 'value

(* | x1 | x2 |
   -----
   | x3 | x4 | *)
type qtree<'value> =
  | Node of qtree<'value> * qtree<'value> * qtree<'value> * qtree<'value>
  | Leaf of treeValue<'value>

[<Struct>]
type Storage<'value> =
  // Storage is always size-x-size square.
  val size: uint64<storageSize>
  val data: qtree<'value>
  new(_size, _data) = { size = _size; data = _data }

[<Struct>]
type SparseMatrix<'value> =
  val nrow: uint64<nrow>
  val ncol: uint64<ncol>
  val nval: uint64<nval>
  val storage: Storage<Option<'value>>

```

Listing 2: Matrix representation

3.1. Basic Data Structures and Functions

Matrix representation, as described in the F# code, is shown in Listing 2. Our representation is generic, meaning it is parameterized by the matrix cell type. The structure does not only work for matrices with sizes that are powers of two. A user-defined matrix of arbitrary size will be aligned to a square storage with size equal to a power of two using Dummy blocks. Since the representation is compressed, the Storage stores the real size (as a power of two) together with the respective tree. Finally, the SparseMatrix stores the matrix dimensions and the number of non-zero entries (nnz) together with data represented as Storage. Data structures for a sparse vector are described in a similar way.

We propose to use Option<'T> to explicitly represent cells that have no values. This approach is natural for graph representation: a cell of the adjacency matrix of a graph with edge labels of type 'T has type Option<'T>: either there exists an edge with a label of the respective type (Some(1)), or there is no edge at all (None). Thus, for a sparse matrix with filled cells of type 'T, the storage has type Option<'T>.

We implemented a minimal set of operations over matrices and vectors required to implement selected graph analysis algorithms. The list of basic functions includes map, fold, reduce. Element-wise operations are based on map2. Particularly, mask is implemented as a case of map2. Linear algebra functions include: vxm (generic vector-matrix multiplication) and mxm (generic matrix-matrix multiplication). We also support transpose and extracting the triangular part of a matrix (required for the triangle counting algorithm). Conversion to and from coordinate list format is provided to simplify input graph building. All operations take sparsity into account and preserve compression. We do not implement fused versions or manually parallelized versions of these functions.

Exception handling is unlikely to be possible in interaction nets in the traditional way due to their highly parallel and irregular nature. Moreover, the selected languages (Inpla and Vine) do not provide any specific mechanisms for exception handling. We focus on realistic code, so we should handle not only correct cases but also erroneous ones (e.g., inconsistent matrix or vector sizes). For this, we use the Result type (even in the F# version)⁵ that explicitly represents two cases: Ok<'R> for successfully completed computations and Err<'E> for errors. This way of handling errors is represented in many

⁵We do this to be as close to interaction-net-based implementations as possible. The F# language has traditional exception mechanisms.

languages such as Rust, F#, etc.

3.2. Implemented Graph Analysis Algorithms

We implemented three basic graph analysis algorithms.

- Level-based BFS: for the given graph and start vertex, it returns a sparse vector `result` such that if `result[i] = Some(k)`, then vertex `i` was visited at step `k`. If `result[i] = None`, then `i` is unreachable.
- SSSP: for the given graph and start vertex, it returns a sparse vector `result` such that if `result[i] = Some(k)`, then the shortest path from the start vertex to `i` has length `k`. If `result[i] = None`, then `i` is unreachable. To prevent an infinite loop in the case of negative cycles, we check that the number of iterations should not be greater than the number of vertices in the input graph.
- Sandia triangle counting algorithm [6] that for the given undirected graph returns total number of triangles in it.

All algorithms were implemented using operations over matrices and vectors similarly to how they are described in GraphBLAS-related materials. We provide basic implementations without any optimizations such as the *push-pull* optimization for BFS or SSSP.

3.3. Interaction Nets Related Details

The translation from F# to Vine and Inpla is almost “as-is” without any interaction-nets-specific optimizations, code practices, or techniques.

In the case of Inpla, we translate high-level programming language functions and algebraic type constructors into an interaction-net agent with a single port and an associated rewriting rule, as shown in Listing 4. Calling such a function involves connecting the function agent to its argument (which may be represented as a tuple agent). The result is then obtained on the agent’s port after a sequence of rewritings. This defines an interface function that is typically rewritten as a single implementation agent with more complex interaction rules.

```
(* Simple interaction rule; real implementation omitted *)
SMap2(r) >< (f, m1, m2) =>
  SMap2Decons_1(r, f, m2) ~ m1;
(* ... *)
(* Operations on generic data - also agents and rules *)
op_add(r) >< (x, y) =>
  op_add_1(r, x) ~ y;

(* Calling the interface function with real data *)
m1 ~ SparseMatrix(4, 4, 16, 4, QTreeLeaf(UserValue(Some(1)))),
m2 ~ SparseMatrix(..., QTreeNode(...)),
SMap2(r) ~ (%op_add, m1, m2);
r; free r;
```

Listing 4: An example of top-level functions and constructors translation result

Inpla has no I/O operations, so input data structures must be generated in advance. Error handling is minimal and relies on the virtual machine’s “no interaction rule” exception for unforeseen situations, as implementing comprehensive checks would introduce significant cognitive overhead in such a low-level language.

Since we do not use any .NET-specific or F#-specific features in our reference implementation, translation to Vine is rather straightforward: this language provides algebraic data types, pattern

matching, and standard `List`, `Option`, and `Result` types. We used the last for error handling. Vine supports file reading, so data loading was implemented. However, this makes performance comparison less accurate because, due to the nature of interaction nets, separation of steps is difficult.

In both Vine and Inpla implementations, weights in SSSP are represented as integers, not floats because Inpla does not support floats, and Vine supports floats but without comparison operators (possibly due to issues with NaN comparison).

4. Evaluation

The goal of this evaluation is to assess the applicability of runtimes based on Interaction Nets to real-world problems, identify bottlenecks and issues, and highlight directions for further investigation.

4.1. Experimental Setup

Hardware. All experiments were conducted on a system with an AMD Ryzen 7 5800X 8-core processor (16 threads with hyper-threading enabled, base clock 3.8 GHz) and 128 GiB of DDR5 RAM.

Competitors. We compare five implementations presented below.

- **F#:** Our quadtree-based single-threaded implementation⁶ that written in F# and compiled for the .NET 10.0 runtime. This serves as a baseline representing a conventional functional language approach and prototype for porting to Interaction Nets. We use BenchmarkDotNet⁷ for benchmarking automation (e.g., for JIT warmup).
- **Inpla (version 0.13.4, patched)**⁸: Our interaction-net-based implementation⁹ in Inpla low-level language, executed on the Inpla runtime system. Our patches of Inpla allow one to use arbitrary heap size configuration with per-thread adjustments via scripting. The patch also increases several virtual machine compile-time constants to support very large trees (thus, large graphs). Additionally, a bottleneck in lists parsing was discovered during testing, and the fix has since been merged into the main Inpla repository¹⁰. Three or five runs for each configuration depending on data size.
- **Vine (commit #41e5ec67)**¹¹: Our interaction-net-based implementation¹² in Vine high-level language, executed on the Vine runtime system. We set 110 GB of heap per worker. Five runs for each configuration
- **LAGraph (version 1.2.1)**: A collection of graph algorithms built on SuiteSparse:GraphBLAS (version 10.3.1), providing state-of-the-art, linear-algebra-based reference implementations of BFS, SSSP, and triangle counting. We use the provided algorithms as a SOTA baseline for high-performance parallel graph analysis. Additionally, we use LAGraph as a baseline to analyze scaling across available threads. We perform 10 runs for each configuration.
- **NetworkX (version 3.6.1)**: A popular Python library for graph analysis. We use the algorithms implemented in NetworkX as a well-known baseline. We do not use any special backends for NetworkX or special settings. We perform 10 runs for each configuration.

⁶Quad-tree-based linear algebra in F#: <https://github.com/Lamagraph/QTreeFSharp>.

⁷BenchmarkDotNet project: <https://github.com/dotnet/BenchmarkDotNet>

⁸Patched version of Inpla: <https://github.com/Lamagraph/inpla>.

⁹Quad-tree-based linear algebra in Inpla: <https://github.com/Lamagraph/QTreeInpla>.

¹⁰Change request to Inpla with lists paring fix: <https://github.com/inpla/inpla/pull/4>

¹¹Installed using instructions from documentation.

¹²Quad-tree-based linear algebra in Vine: <https://github.com/Lamagraph/QTreeVine>.




Figure 4: Matrices shapes: original and after reordering

Dataset. We evaluate our implementations on several real-world and synthetic graphs from the SuiteSparse Matrix Collection [7]. The graphs are chosen to represent a variety of sizes, sparsity patterns, and application domains, taking into account the capabilities and limitations of the selected competitors. Table 1 summarizes the number of vertices ($|V|$) and edges ($|E|$) of the selected graphs. Weights were shifted relative to the minimum weight per graph to prevent negative cycles and were converted to `Int` because `Inpla` does not support floats. All graphs were converted to undirected and have exactly one connected component; thus, any vertex can be selected as a start for BFS and SSSP. We selected Vertex 0 for all runs.

Table 1: Dataset properties

Graph	$ V $	$ E $
coAuthorsCiteSeer	227,320	1,628,268
boneS01	127,224	5,516,602
usroads-48	126,146	323,900
Andrews	60,000	760,154
Trefethen_2000	2000	41,906
gr_30_30	900	7,744

Reordering. Sparse matrix reordering techniques can improve performance of operations over matrices [8]. In case of quad-tree based representation reordering also can reduce memory consumption. We use reverse Cuthill–McKee for reordering. Patterns for selected matrices is represented in Figure 4. Patterns for *gr_30_30* is similar to patterns for *boneS01*. In evaluation we use both original and reordered matrices except *coAuthorsCiteSeer* that used reordered only version because original version cause OOM of `Inpla`. We can see, that in this case reordered matrix has big empty blocks that represented in quad-tree compactly. Similarly for *Andrews*. For other matrices reordering effect was investigated in evaluation.

Research Questions. Our evaluation addresses the following research questions:

RQ1 (Performance). How does the execution time of `Inpla` and `Vine` compare to `LAGraph`, `NetworkX` and `F#` on BFS, SSSP, and triangle counting?

RQ2 (Scaling). How well do the `Inpla`, `Vine`, and `LAGraph` implementations scale with the number of available CPU cores? We measure speedup from 1 to all available cores for BFS, SSSP, and triangle counting.

RQ3 (Impact of Matrix Reordering). How does reordering affect the performance of competitors?




Figure 5: Evaluation results for BFS on *coAuthorsCiteseer* (reordered)




Figure 6: Evaluation results for BFS on *boneS01* and *usroads-48* (both original and reordered)




Figure 7: SSSP evaluation results for *Andrews* (both original and reordered)




Figure 8: Evaluation results for *Trefethen_2000* (both original and reordered)




Figure 9: Evaluation results for *gr_30_30* (both original and reordered)




Figure 10: Additional scaling analysis

4.2. Evaluation Results

First, we discuss the limitations that prevented evaluating all (graph, algorithm, tool) combinations. The primary limitation is the substantial memory consumption of Inpla and Vine. Since our implementations rely on a recursive divide-and-conquer computation pattern, aggressive duplication occurs during computations. While BFS and SSSP require only the *vxm* operation, triangle counting requires the *mxm* operation, which involves more complex recursive calls. As a result, triangle counting is feasible only on relatively small graphs. Notably, Vine encounters out-of-memory (OOM) errors even on *Trefethen_2000* when performing triangle counting. Second, Vine is excessively slow. Therefore, we can execute it within a reasonable time only on small matrices, namely *Trefethen_2000* and *gr_30_30*. Additionally, Inpla has dataraces that leads to instability¹³. We check that results are correct, but des not guarantee that raises does not affect performance patterns.

Regarding **performance** we can see, that in all cases (Figures 5a, 6a, 7a, 8a, 8b, 9a, 9b) Interaction Nts based implamantations in several orders of magnitude slower (up to 10^5) then both NetworkX and LaGraph. Note that in almost the all cases LaGraph is faster that NetworkX. So for all selected cases linear algebra based approach can demonstrate reasonable performnnc. We can conclude that if graph is big enough (*coAuthorsCiteseer*, *boneS01*, *usroads-48*, *Andrews*) then Inpla can outperforms F# implementation. Note that

Inpla slower in case SSSP on *gr_30_30* Moreover inpla scales very bad in this case. It may be cased by manual tuning of parameters aimed on performance on big grpahs.

Scaling. Near to linear scaling on up to 4 threads. Inpla scales better than LaGraph: Complex scaling pattern of LaGraph and more evident for Inpla and Vine. Vine scale better than Inpla. On small data (*Treferthen_2000* and *gr_30_30*) Inpla scales poor.

Reordering analysis shows that in cases when initial matrix has no specific pattern (*coAuthorsCiteseer* and *Andrews*) it helps to reduce memory consumption and to improve performance. In cases when

¹³Some of them already fixed by our team (e.g. this one: <https://github.com/inpla/inpla/pull/8>), other must be fixed in the future.

matrix has diagonal-like initial pattern contribution of reordering is not so trivial. Inpla: For *usroads-48* reordered faster. For *bone501* reordered is slower. Note that for these graphs F# demonstrates no difference in performance between original and reordered matrices. For Vine for almost all cases reordered slower. In some cases reordering does not affect performance significantly. While reordered slower, it scales better.

Also we analyze total number of reductions presented in Figure 11. As expected, Vine performs significantly more reductions because compile code to fixed set of basic agents while Inpla allows user to use custom agents. Inpla way can be better to optimize domain-specific computations.

To summarize.

- Coding patterns together with evaluation strategy=y and memory management mechanisms to reduce memory consumption.
- Load balancing (reordering and scaling)
- Configuration of Inpla (performance on small graphs)
- Anomaly number of reductions and performance for SSSP on *gr_30_30* Gap in number of reductions is smaller than in other cases. Performance of Vine is better than Inpla. SSSP only. BFS ok.




Figure 11: Total number of interactions

A significant concern is memory consumption. Interaction nets programming requires the use of duplicator agents for data reuse. Inpla's current implementation eagerly copies entire subnets when duplicating, resulting in substantial memory overhead compared to the F# prototype. Vine does not, but still requires a huge amount of memory. Addressing this issue requires further investigation.

5. Conclusion and Future Work

We have demonstrated a feasible mapping of GraphBLAS concepts to interaction nets using quadrees, provided concrete implementations in Inpla and Vine, and evaluated their scalability. The performance gap compared to state-of-the-art parallel linear-algebra-based frameworks remains substantial, and scalability is limited.

Languages: But for comfortable usage Result monad or similar mechanisms is required.

Several bottlenecks and limitations were identified. The naive implementation suffers from a duplication problem: a huge number of duplicated nodes limits both performance and memory efficiency. Reordering (RCM) is crucial for performance and memory consumption but is not handled automatically. Currently it is a separated preprocessing step. Moreover, only a limited set of operations and algorithms has been evaluated.

Future work includes investigating improved duplication strategies. The set of implemented operations and algorithms should be extended to include indexed operations, slices, and such graph analysis algorithms as PageRank, minimum spanning forest (MSF), regular path queries, context-free path queries (CFPQ), and maximal independent set (MIS).

A systematic evaluation of other interaction-net-based implementations, such as HVM, is also required. Alternative sparse matrix representations, such as the coordinate list (COO) format, rely on sorting—an operation that parallelizes well in interaction nets and therefore warrants further investigation.

Classical implementations of algorithms. May interaction nets be good enough to parallel straight-forward implementations?

Automatic kernel fusion presents another direction. Advanced optimization techniques natural to functional programming, such as supercompilation [9] or distillation [10], could be applied to

automatically fuse sequences of operations, analogous to manual fusion for masks and other constructs in classic GraphBLAS. Partial evaluation for interaction nets [11] Whether these techniques can be adapted for interaction nets remains an open question.

Finally, hardware acceleration remains an open opportunity. Designing an FPGA specifically optimized for interaction net reduction could exploit the fine-grained parallelism inherent to the model and yield substantial performance gains.

Acknowledgments

This research has been supported by the St. Petersburg State University, grant id 116636233.

Declaration on Generative AI

During the preparation of this work, the authors used !!!! in order to: Grammar and spelling check. After using these tool, the authors reviewed and edited the content as needed and take full responsibility for the publication's content.

References

- [1] B. Brock, A. Buluç, T. G. Mattson, S. McMillan, J. E. Moreira, Introduction to graphblas 2.0, in: 2021 IEEE International Parallel and Distributed Processing Symposium Workshops (IPDPSW), 2021, pp. 253–262. doi:10.1109/IPDPSW52791.2021.00047.
- [2] T. Mattson, T. A. Davis, M. Kumar, A. Buluc, S. McMillan, J. Moreira, C. Yang, Lagraph: A community effort to collect graph algorithms built on top of the graphblas, in: 2019 IEEE International Parallel and Distributed Processing Symposium Workshops (IPDPSW), 2019, pp. 276–284. doi:10.1109/IPDPSW.2019.00053.
- [3] Y. Lafont, Interaction nets, in: Proceedings of the 17th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL '90, Association for Computing Machinery, New York, NY, USA, 1989, p. 95–108. URL: <https://doi.org/10.1145/96709.96718>. doi:10.1145/96709.96718.
- [4] N. Huber, W. Yi, An encoding of interaction nets in ocaml, Electronic Proceedings in Theoretical Computer Science 417 (2025) 1–16. URL: <http://dx.doi.org/10.4204/EPTCS.417.1>. doi:10.4204/eptcs.417.1.
- [5] I. Mackie, S. Sato, In-place graph rewriting with interaction nets, in: A. Corradini, H. Zantema (Eds.), Proceedings 9th International Workshop on Computing with Terms and Graphs, TERMGRAPH 2016, Eindhoven, The Netherlands, April 8, 2016, EPTCS, 2016, pp. 15–24. URL: <https://doi.org/10.4204/EPTCS.225.4>. doi:10.4204/EPTCS.225.4.
- [6] M. Wolf, M. Deveci, J. W. Berry, S. D. Hammond, S. Rajamanickam, KKTri: Fast Linear Algebra-Based Triangle Counting with KokkosKernels., Technical Report, Sandia National Laboratories (SNL-NM), Albuquerque, NM (United States), 2017.
- [7] S. P. Kolodziej, M. Aznavah, M. Bullock, J. David, T. A. Davis, M. Henderson, Y. Hu, R. Sandstrom, The suitesparse matrix collection website interface, Journal of Open Source Software 4 (2019) 1244. URL: <https://doi.org/10.21105/joss.01244>. doi:10.21105/joss.01244.
- [8] O. Asudeh, S. Mahdipour Saravani, F. Rastello, G. Sabin, P. Sadayappan, How effective is matrix reordering for improving performance of sparse matrix-vector multiplication?, in: Proceedings of the SC '25 Workshops of the International Conference for High Performance Computing, Networking, Storage and Analysis, SC Workshops '25, Association for Computing Machinery, New York, NY, USA, 2025, p. 823–827. URL: <https://doi.org/10.1145/3731599.3767441>. doi:10.1145/3731599.3767441.
- [9] V. F. Turchin, Supercompilation: Techniques and results, in: D. Bjørner, M. Broy, I. V. Pottosin (Eds.), Perspectives of System Informatics, Springer Berlin Heidelberg, Berlin, Heidelberg, 1996, pp. 227–248.

- [10] G. Hamilton, The next 700 program transformers, in: E. De Angelis, W. Vanhoof (Eds.), *Logic-Based Program Synthesis and Transformation*, Springer International Publishing, Cham, 2022, pp. 113–134.
- [11] D. Béchet, Partial evaluation of interaction nets, in: M. Billaud, P. Castéran, M. Corsini, K. Musumbu, A. Rauzy (Eds.), *Actes WSA'92 Workshop on Static Analysis* (Bordeaux, France), September 1992, Laboratoire Bordelais de Recherche en Informatique (LaBRI), Proceedings, Series Bigre, Atelier Irisa, IRISA, Campus de Beaulieu, 1992, pp. 331–338.